

Mounika Raj

+1 (479) 640-0751 | mounikacareer123@gmail | LinkedIn

WORK EXPERIENCE

NVIDIA Corporation

TX, USA

ML Engineer

Feb 2026 - Present

- Architected distributed ML inference infrastructure supporting deployment of perception and prediction models across GPU clusters, reducing end-to-end inference latency by 34%
- Designed CI/CD automation for ML workflows using GitHub Actions, unit tests, and rollout policies, improving deployment reliability from 83% → 99.4%.
- Optimized multi-GPU training and inference workflows using CUDA-enabled environments, improving accelerator utilization by 41% and reducing model execution time by 38%.
- Developed model export and deployment pipelines utilizing ONNX and TensorRT, reducing inference overhead by 33% across production GPU workloads.
- Built automated validation frameworks for ML model deployment, identifying numerical regressions and improving inference reliability by 29% across production releases.

Meta Platforms, Inc.

India

AI Engineer

Oct 2021 - Nov 2023

- Engineered scalable ML platform services supporting distributed model training and inference workloads processing 32M+ requests monthly.
- Built large language model pipelines processing 32M+ conversational interactions, reducing response latency by 34% within AI-driven systems.
- Designed Kubernetes-based ML orchestration workflows enabling reproducible model training, deployment, and monitoring across cloud environments.
- Developed high-throughput data processing pipelines integrating distributed storage and compute infrastructure, increasing ML pipeline throughput by 37%.
- Built cloud-native ML infrastructure utilizing Kubernetes, Docker, and Python services supporting large-scale AI experimentation.
- Developed model lineage and metadata tracking solutions improving reproducibility and deployment governance across enterprise ML platforms.

JPMorgan Chase Pvt. Ltd

India

ML Engineer

Oct 2019 - Sep 2021

- Architected distributed data processing systems managing 14M+ records daily, improving platform scalability and operational throughput by 35%.
- Developed high-performance microservices using Python and C++ supporting real-time analytics and distributed inference workloads.
- Coordinated development and production rollout of machine learning risk assessment frameworks, ensuring model reliability, compliance alignment, and scalable financial analytics operations.
- Optimized data analytics pipelines processing 5.1TB+ transactional datasets monthly, accelerating reporting and model evaluation throughput by 38%.
- Architected predictive modeling solutions integrating SQL, Python, and distributed processing environments, improving operational decision-making efficiency by 21%.
- Implemented real-time ML pipelines for sensor data analysis, reducing inference latency by 30% through model compression and GPU optimization.

Intel Technology Pvt. Ltd.

India

AI Engineer

Jun 2017 - May 2019

- Developed computer vision inference pipelines processing semiconductor inspection data using GPU-accelerated deep learning architectures.
- Engineered deep learning models processing 9M+ semiconductor sensor records, increasing anomaly detection accuracy by 31% across validation systems.
- Optimized computer vision algorithms for automated defect inspection workflows, reducing false-positive identification rates by 23%.
- Built distributed ML training infrastructure utilizing TensorFlow, Kubernetes, and GPU clusters, reducing training duration by 42%.
- Developed NLP-based document classification systems with BERT and spaCy to automate data labeling, reducing manual processing time by 50%.
- Implemented real-time recommendation system enhancements using graph neural networks (GNNs), leading to a 10% increase in user session duration across pilot markets.

EDUCATION

Webster University
Master in Cybersecurity

MO, USA
Graduation Date: Dec 2025

Bharat Degree College For Women
Bachelor in Computer Science

India
Graduation Date: Apr 2017

SKILLS :

Programming Languages	Python, SQL, R, C++, Java
Machine Learning & AI	Supervised/Unsupervised Learning, Deep Learning, Computer Vision, Natural Language Processing (NLP), Generative AI, Large Language Models (LLMs), Reinforcement Learning
Frameworks & Libraries	TensorFlow, PyTorch, Keras, Scikit-learn, Hugging Face Transformers, LangChain, OpenCV, XGBoost, spaCy
MLOps & Deployment	MLflow, Airflow, Docker, Kubernetes, AWS SageMaker, Azure ML, Databricks, CI/CD Pipelines
Data Engineering & Visualization	Pandas, NumPy, Spark, SQL, Tableau, Power BI, Matplotlib, Seaborn
Cloud & Tools	AWS EC2, AWS S3, GCP Compute Engine, Azure DevOps, GitHub Actions, GitLab, Terraform, Linux/Unix, Bash Scripting, Jira
Soft Skills & Collaboration	Cross-Functional Communication, Problem Solving, Agile Methodologies, Leadership, Technical Documentation